

Mock Interview — Instructions

1. Why we do this

Most of you will sit a real interview within 3 months — for an internship, a graduate position, or a master’s programme. Interview performance is a skill that improves quickly with practice and feedback, and almost not at all with reading about it. Two rounds, one at the start of the semester and one at the end, let you measure your own improvement against a fixed standard.

You will be on both sides of the table. Sitting on a panel (i.e., interviewer) teaches you more about what a weak answer sounds like than being asked the question does, and that knowledge transfers directly to your own answers.

2. Structure at a glance

	Round 1 — 17 September	Round 2 — 17 December
Purpose	Baseline diagnostic, identify weaknesses early	Measure improvement, rehearse for Oct. – Dec. applications
Format	Individual interview before a panel of 3–4	Individual interview before a panel of 3–4, upgraded content
Pairing	G1–G2, G3–G4, G5–G6	G1–G4, G2–G6, G3–G5 (no repeats)
Roles	Both directions within the session — every group interviews and is interviewed	Same
Content	Motivation, behavioral, light technical, estimation, AI use	The above plus your project findings, method defence, and programme-fit questions
Questions	Supplied in this document — no preparation required	Drawn from the Track 3 question bank plus the interviewee’s own presentation

Three panels run concurrently, one per paired set of groups. Each session has two halves: in the first, one group of each pair sits as the panel; in the second, the roles reverse. Everybody is interviewed once and sits on a panel once, in every round.

3. Round 1 — Thursday 17 September

Running order

Time	Activity
13:00–13:08	Briefing; panels take positions; role cards and feedback forms distributed
13:08–13:38	Half A — Groups 1, 3, 5 as panels; Groups 2, 4, 6 interviewed. Six minutes per candidate, then written feedback
13:38–13:41	Switch
13:41–14:11	Half B — roles reversed
14:11–14:15	Plenary debrief; forms collected

If you are on the panel

- Decide among yourselves who asks which question before the first candidate sits down. Do not improvise the order.
- Ask three questions from the pool below — one motivation, one behavioural, one technical or estimation — then one follow-up on whichever answer was weakest.
- **Ask the follow-up.** This is the part panels skip and the part that matters: “You said the data were messy — messy how?” or “What would you have done if that had not worked?” An unchallenged answer teaches the candidate nothing.
- Complete one feedback form per candidate. Hand it to them directly at the end of the half-round.
- Stay in role. No coaching mid-interview, no reassurance, no laughing along.

If you are the candidate

- Six minutes is short. Lead with your answer and then support it; do not build up to the point.
- Specifics beat adjectives. “I cleaned a 4,000-row sales file with inconsistent date formats” is an answer; “I have strong data skills” is not.
- If you do not know something, say so and say what you would do to find out. Bluffing is the single most common failure in this room and the easiest to detect.
- **Log the questions you were asked** in the shared class sheet immediately afterwards, with a one-word difficulty rating. This becomes baseline data for the Track 3 team.

Question pool — Round 1

Panels choose from this pool. Candidates should not rehearse specific answers; the point is to find out where you are.

Motivation and fit. Tell me about yourself in ninety seconds. Why analytics rather than finance or marketing? What do you want to be doing three years from now, and what makes you think you would be good at it?

Behavioural. Describe a time you worked with data that turned out to be worse than you expected. Tell me about a disagreement in a team project and how it resolved. Describe something you taught yourself outside class.

Technical. Explain a p-value to a marketing manager who has never taken statistics. What is the difference between correlation and causation, and give me a business example where confusing them would be expensive? You receive a customer transactions file — what are the first four checks you run? Explain what a join does, without using the word “join”.

Estimation. How many bubble tea shops are there in Wenzhou? How much does this university spend on electricity in a year? Estimate the number of BA-related internship positions advertised in Shanghai this month.

AI use. How do you use AI tools in analytical work, and what do you check before you trust the

output? Tell me about a time a model gave you something that was wrong.

4. Round 2 — Thursday 17 December

The same structure, with harder content and a re-pairing so that nobody faces the same panel twice. By this date you will have presented your project (10 or 15 December), so panels will have seen the work of the group they are interviewing and can ask about it.

What changes

- **Seven minutes per candidate** instead of six.
- **Questions come from the Track 3 question bank**, compiled from real interview accounts — not from the pool above.
- **Two questions must concern the candidate's own project:** walk me through how you obtained your data; what failed and what you did about it; which of your findings are you least confident in, and why; defend the number in your headline chart.
- **One question is programme-fit**, for those applying to master's programmes: why this programme, what will you do with it, what would you want to study in more depth.
- **Panels compare against the Round 1 form.** Bring your own Round 1 feedback with you; the closing debrief is about what changed.

5. Feedback and assessment

Both rounds count towards Participation. Round 1 is assessed on attendance, staying in role, and the quality of the feedback you give — not on how you perform as a candidate. Round 2 adds a light performance element.

Panels score each candidate on four criteria. A single number is useless; the written comments are the part candidates actually use.

Criterion	4 — Ready	3 — Solid	2 — Developing	1 — Not yet
Structure	Point first, then support; finishes inside the time	Clear but takes a while to arrive	Rambles; the listener assembles the point	No discernible shape
Evidence	Named tools, numbers, concrete situations	Some specifics, some generalities	Mostly adjectives	Claims with nothing behind them
Technical accuracy	Correct, and pitched to the listener	Correct but jargon-heavy or over-simplified	Partly wrong or evasive	Wrong, or bluffed
Under follow-up	Engages, revises, admits limits	Holds up with effort	Repeats the first answer louder	Collapses or bluffs

Every form ends with two free-text lines: one thing this candidate should keep doing, and the single change that would most improve the next answer. Write the second one as an instruction, not an

observation — “give the number before the story” rather than “was a bit unstructured”.

6. What to submit

- 1) **The feedback forms you completed as a panel member**, uploaded to Canvas by the Sunday following each round.
- 2) **Your question log** — the questions you were asked, entered in the shared class sheet the same day.
- 3) **A reflection of no more than 100 words**: after Round 1, the one thing you will work on before December; after Round 2, whether it changed and what the evidence is.

7. Preparation

For Round 1, none. Come as you are. A rehearsed baseline is not a baseline, and it will cost you the diagnostic value of the session. Read the question pool once so you know the shape of the thing; do not script answers.

For Round 2, prepare properly. Know your project well enough to be challenged on it, re-read your Round 1 feedback, and have a ninety-second account of your project that a non-specialist would follow.

The one rule that matters in both rounds: take it seriously for the sixty minutes it lasts. A panel that laughs along, or a candidate who performs ironically, wastes the session for the three other people in the group. You will each get roughly six minutes of interview practice in September. There is no spare.

Questions in class or at office hours: Tue 16:00–17:15, Wed 15:00–16:00, Thu 14:15–17:00, CBPM B214.

How AI was used to create this instruction

This course asks you to disclose your use of AI, so here is mine. I drafted these instructions with Claude (Anthropic, Opus 5 Extra): structuring the two rounds, drafting and redrafting prose, and checking the timings against the class schedule. The design decisions are mine — the paired-group format, reversing roles within each session rather than across the two rounds, the choice of panel over one-to-one interviews, the decision not to grade Round 1 on performance, and the question pool. Several proposed structures were rejected, and the arithmetic in the running order was checked by hand rather than taken on trust. If you find that something here does not work in the room on Thursday, tell me; that is useful information, not a complaint.